DIAGNOSTIC ARTIFACT

Outlier AI — Structural Diagnosis of Recurring Platform Failure

Prepared by: Jamie Forrester, Independent Systems Architect
Engagement type: Public Diagnostic Exemplar
Date: April 2026
Classification: Public diagnostic exemplar — derived from a full internal diagnostic artifact and based entirely on open sources

---

Evidentiary standard: This artifact distinguishes directly evidenced facts [A][B], convergent public symptom patterns [C], consequence evidence [D], and structural inference [E]. The governing diagnosis is a high-confidence structural inference from convergent public evidence — not a claim about internal intent. Evidence tier references appear inline; source classes are defined in Section 14.

Judgment boundary: This artifact makes no claim about intentional deception, fraud, or bad faith by Outlier or Scale AI. It is a structural reading of public evidence. The mismatch described is consistent with rapid scaling under conflicting commercial pressures. No inference is made about individual intent, internal dashboard logic, undisclosed client instructions, or adjudicated legal liability beyond what public sources directly support.

SECTION 0 — WHAT THIS DOCUMENT IS

This is a structural diagnosis from public evidence.

It is not a legal brief, an operational audit, a moral judgment, a technology assessment, or a redesign specification. It identifies the governing primitive, the first structural break, the function-collapse condition, the compensatory burden layer, and the high-level dependency order for correction — derived entirely from publicly observable system behaviour.

This public version preserves the core diagnostic spine to demonstrate method, but does not disclose the complete replacement specification. The full internal redesign artifact specifies the target-state architecture, transition sequence, and implementation logic. That is a separate object.

SECTION 1 — EXECUTIVE SUMMARY

Outlier has a recurring platform failure that its current correction loops cannot resolve. Credentialed contributors fail production. Quality scores do not predict output quality. Qualified contributors sit idle in Empty Queue states. High-signal specialists churn fastest while reporting the work is not what they were recruited to do. Successive interventions — stricter credentials, longer guidelines, Marketplace visibility, dashboard redesigns, Queue Manager support, webinars — have not reduced the core friction. They have only moved it.

The governing diagnosis: Outlier recruits through the grammar of an expertise market while operating through the mechanics of a reliability pipeline.

An expertise market purchases open-ended specialist judgment and values correct variance. A reliability pipeline purchases repeatable instruction adherence under constraint and suppresses variance through qualification gates, unitised tasks, rubrics, and continuous quality control. Outlier's recruitment surface signals the first. Its execution layer runs the second. Because the intake architecture and the execution architecture are selecting for different inputs, every downstream mechanism — the qualification gate, the quality score, the routing algorithm, and the payment structure — operates on a false premise.

The first structural break occurs at the qualification-to-execution handoff. Contributors are selected through expertise-coded signals — degrees, LinkedIn profiles, domain screenings, identity verification — then moved into success conditions dominated by rubric adherence, timing, formatting, and bounded outputs. The work shifts from *use your expertise* to *match the manual*. The expertise is not consumed. Compliance is consumed. The strongest structural reading is that the architecture derives legitimacy benefits from expert-coded recruitment while operationally consuming compliance-governed output.

This is why every prior fix has failed. Stricter credentials tighten a gate already measuring the wrong thing. Longer guidelines intensify the signal-switch, converting expert labour into compliance labour earlier and more completely. Each intervention addresses its own layer. None touch the primitive. Until the primitive is corrected, future fixes will reproduce the same failure in new forms.

A second-order consequence is the emergence of compensatory human labour. Queue Managers — contributors, not employees — function as human bridges between opaque routing, shifting guidelines, and contributor confusion. Their role is load-bearing, not supplementary.

What must change is not the interface. What must change is the primitive. If Outlier is fundamentally a reliability pipeline, it must recruit, qualify, route, measure, and compensate for reliability work. If it also requires expert judgment, that work must be separated into a distinct stream with different gates, different success criteria, and an upstream signal channel.

One-line verdict for the board: Outlier is an expertise-coded reliability pipeline. Its recurring failure is not a contributor problem. It is a category error embedded in the system's primitive. Correct the primitive, or every future fix will fail in the same way.

SECTION 2 — ENGAGEMENT SCOPE

System under analysis
The Outlier contributor platform — the full contributor lifecycle from recruitment surface through qualification, onboarding, task routing, quality-scoring, queue-state, support layers, payment, and the external consequence layer where legal and reputational pressure have appeared.

An important framing note
The structural mismatch diagnosed here is not a moral failure. It is the predictable output of a system scaled rapidly under commercial pressure — Scale AI's obligation to deliver reliable, high-throughput training data to frontier AI clients — while its public recruitment surface borrowed the legitimacy grammar of expert sourcing to attract credentialed contributors and reassure those same clients. The two requirements point in opposite directions. The mismatch is structural, not intentional.

Parent system context
Outlier's governing constraint is not determined by its own recruitment surface but by Scale AI's upstream commercial mandate: reliable, scalable human-generated inputs for AI development clients including OpenAI, Meta, Microsoft, Google, and Anthropic [A]. This parent constraint is treated as fixed.

Evidence basis
This diagnosis is based entirely on publicly available materials. Primary sources [A]: Outlier's official website, FAQ, blog posts, specialist job pages, Marketplace announcements, Community Guidelines, and Terms of Service; Scale AI's official materials; the Oxford Economics commissioned report (December 2025). Symptom evidence [C]: Glassdoor, Indeed, Trustpilot, Quora, Reddit (r/outlier_ai), BBB complaints, and a community-built GitHub browser extension (EmptyQueue-Extension). Consequence evidence [D]: McKinney v. Scale AI (December 2024), Amber Rogowicz PAGA suit (January 2025), Bloomberg Law WARN Act reporting (October 2024), TechCrunch DOL investigation reporting (March 2025), The Guardian (April 2026), Oxford Internet Institute labour standards assessment (2023).

Known fixed constraints: contractor classification model; client NDA and codename structure; demand-driven project availability; algorithmic quality scoring backbone; core task types (rating/ranking, prompt creation, editing, annotation).

SECTION 3 — FAILURE PATTERN AS PRESENTED

The following synthesises how platform leadership, operations, and contributors describe the problem — in their own language, before structural translation:

"We recruit highly credentialed, domain-specific contributors through expertise-coded signals, but contributor success at intake does not reliably predict success inside the production workflow." [A][C]

"High-signal contributors either fail early assessment gates, enter production and encounter work materially different from what they expected, or churn rapidly." [C]

"Quality variance persists despite increasingly strict compliance layers." [A][C]

"The Empty Queue is endemic. Contributors qualified for specific work cannot access that work." [C]

"Support functions — Queue Managers, webinars, office hours, war rooms — are becoming permanent compensatory infrastructure." [C]

These are not isolated complaints. They describe recurring friction across qualification, production, routing, and retention [A][C]. The pattern is identical at every lifecycle stage — the first structural signal that this is not an execution problem.

SECTION 4 — EARLY STRUCTURAL SIGNALS

Four structural signals are visible from the public record before the formal diagnosis is applied:

The two grammars diverge at the handoff.
The recruitment surface is uniformly expertise-coded: degrees, LinkedIn profiles, academic verification, specialist domain selection [A]. The execution surface is uniformly reliability-coded: standardised tasks, time limits, 50+ page rubrics, algorithmic scoring by formatting compliance and time-on-task [A][B]. These are two different grammars describing two different systems, operating on the same contributor population with no reconciliation mechanism between them.

Every correction has been interface-level.
Every public response to the recurring failure has addressed a surface: routing interface, navigation interface, compliance interface, intake filter [A]. Because the system has not explicitly declared its true production primitive, correction efforts remain confined to interface-level adjustments rather than primitive-level redesign.

Compensatory human structures have become load-bearing.
Queue Managers exist to absorb system ambiguity that the platform's own information

architecture cannot carry [C]. When an organisation requires a permanent human layer to interpret its own operational conditions to the people doing its core work, the system's design has failed at a level below the interface.

The external pressure is patterned, not random.
Legal action, labour complaints, investigative journalism, and public symptom clustering all point to the same tension: the gap between what the platform signals it is and what contributors experience once inside it [C][D]. Patterned external pressure is a consequence signal, not a cause.

SECTION 5 — SYSTEM BOUNDARY MAP

— Figure 1: Category Divergence Map —

The platform's structural divergence can be described as two incompatible grammars operating simultaneously on the same contributor population, with the qualification gate as the point where the mismatch becomes operational:

Feature	Recruitment grammar (expertise-coded)	Execution grammar (reliability-coded)
What is signalled	Specialist judgment, credentialed access, expert contribution	Rubric adherence, task completion, compliance
Success metric	Depth of domain expertise	Formatting, time-on-task, peer reviewer agreement
Feedback loop	Specialist validation	Algorithmic scoring, EQ state
Unit of work	Open-ended problem, judgment call	Bounded task with defined output format
Variance	Valued — correct variance is the product	Suppressed — variance is the failure signal
Memory	Implied portfolio, cumulative expertise	Per-project only, no portable record

Queue Managers appear as a compensatory loop outside the main pipeline, absorbing the ambiguity generated by the grammar mismatch. The missing Upstream Signal Channel is a structural absence — no interface exists for contributors to surface reframing insight to a layer where it can be acted on.

Primary actors:

Actor	Role	Position in system
Prospective contributor	Self-selects based on recruitment grammar	Pre-intake
Contributor (Attempter)	Completes entry-level tasks under rubric	Production layer
Contributor (Reviewer)	Quality-assesses other contributors' work	QA layer
Contributor (Oracle)	Senior review and escalation	QA escalation
Queue Manager (QM)	Interprets system ambiguity; hosts webinars and war rooms	Coherence layer
Algorithmic scoring system	Detects quality; triggers access, payment, removal	Authority layer
Project routing system	Matches contributors to available tasks	Routing layer
Client (codename-protected)	Sets task definitions, rubrics, project scope	Upstream
Scale AI	Sets commercial mandate, owns client relationships	Parent

Key interfaces:

- Recruitment interface: expertise-coded; sets intake expectation [A]
- Qualification gate: uses expertise signals; scores compliance [A][B]
- Project-specific onboarding: introduces task rubrics; first point of signal-switch [B][C]
- Task execution environment: time-constrained, rubric-governed, algorithmically scored [A][B]
- Queue and routing system: internal EQ reason codes not surfaced to contributors [B][C]
- QM support layer: human interpretation bridge — compensatory, not designed [C]
- Payment system: weekly; subject to algorithmic reduction for time-limit overruns [B][D]

Decision points where structural failure concentrates:

- What constitutes a passing assessment — currently measures wrong input [A][B][C]
- How quality scores are calculated — currently measures compliance, not expertise [B][C]
- Why a contributor's queue is empty — currently withheld by design [B][C]

- Whether a contributor can appeal a score or deactivation — currently: no mechanism [B][D]
- Whether upstream reframing feedback is collected — currently: no such interface exists [A][B]

SECTION 6 — WORKFLOW AND PIPELINE TRACE

The documented operating sequence:

Recruit → Qualify → Verify → Onboard → Constrain → Perform → Score → Route → Repeat or Remove

1. Contributor creates account, uploads resume, selects skills, completes ID verification [B]
2. System cross-references LinkedIn, validates academic email, verifies publications [B]
3. Contributor is offered project-specific onboarding — or hits Empty Queue immediately [A][C]
4. Contributor reads guidelines (50+ pages), may attend webinars or office hours — inconsistently compensated [B][C][D]
5. Contributor completes qualification task graded by AI or peer reviewers [A][B]
6. **If pass:** contributor enters production queue. **If fail:** no explanation provided [C]
7. Contributor receives unitised task with time limit, rubric, and output format [A][B]
8. Response scored algorithmically or by peer reviewers [A][B]
9. Score affects future task access, project eligibility, tier status [B][C]
10. When demand drops or project ends: blank dashboard, no explanation, no timeline [C]
11. Contributor seeks help from Queue Managers — webinars, Slack, war rooms [C]
12. Contributor either continues through cycles or churns [C][D]

Critical structural observations:

There is no step for upstream reframing [A][B]. The workflow accepts output only — no mechanism exists for a contributor to surface a malformed rubric, an unrealistic time limit, or a task defined for the wrong downstream use.

Memory is per-project only [A][B]. A contributor who excelled on Project A cannot carry that signal to Project B. Every project begins assessment from zero.

The Empty Queue is not a workflow step. It is a silent capacity management mechanism [B][C].

Queue Managers are inserted after the fact — a structural patch, not a designed component [C].

SECTION 7 — FUNCTIONAL DIAGNOSIS

Applying the Four-Function Law [E]: institutions fail under scale when sensing, interpretation, authority, and memory remain fused or misaligned at the point of consequence.

Function	Location	Condition	Consequence
Sensing	Algorithmic scoring + peer review	**Mis-specified** — calibrated to detect rubric adherence proxies rather than the recruited input class	Quality detection operates on the wrong signal [A][B][C]
Interpretation	QM layer + guideline documents	**Fragmented** — displaced into informal, non-authoritative human support with no design authority	Meaning of quality varies by project and reviewer [C]
Authority	Automated system	**Fused with sensing** — automated consequence without review separation	Mis-specified sensing triggers consequence at scale [B][D]
Memory	Per-project only	**Local and non-portable** — no cross-project contributor record	System cannot learn; routing cannot improve [A][B][C]

The critical fusion: Sensing and authority are fused in the algorithmic scoring system [B][C][E]. The same mechanism that detects quality also drives deactivation, pay reduction, and access changes — with no human review layer between detection and consequential action. A mis-specified sensor with executive authority, operating across 100,000+ contributors, produces systematic misclassification as a structural output.

The critical absence: Cross-project memory. The system discards its own evidence of contributor capability at the end of every project cycle [B][C].

The critical fragmentation: Queue Managers are the de facto interpretation layer but have no design authority and cannot correct upstream errors [C]. Their interpretations are local, inconsistent, and ungoverned.

SECTION 8 — STRUCTURAL FAILURE POINT

The first structural break: The qualification-to-execution handoff.

This is the moment a contributor who entered through expertise-coded recruitment receives their first project-specific rubric. It is where the platform's two grammars meet and where the mismatch becomes operational.

What happens: The contributor was selected using expertise signals — degree, domain knowledge, specialist identity [A][B]. At the handoff, they receive a 50+ page document specifying exactly how tasks must be completed: formatting requirements, length limits, time constraints, response structures, scoring rubrics [A][B][C]. The question shifts from *what is the correct answer given your expertise?* to *what answer conforms to this rubric?* When they diverge, the scoring layer sides with the rubric every time. The expertise is not consumed. Compliance is consumed.

What the break is not:

- It is not a guideline clarity problem — clearer guidelines intensify the signal-switch [E]
- It is not a contributor skill problem — the system is not primarily designed to consume open-ended specialist judgment as its dominant production input [A][B][E]
- It is not a routing inefficiency — routing is downstream of the handoff [E]

The three capability modes — how the system handles each:

Mode	Description	How Outlier handles it
Execution reliability	Stable instruction adherence under monitoring	Actively selected for. This is the actual operating need. [A][B]
Diagnostic capability	Error detection within defined criteria	Partially present via Reviewer and Oracle tiers. [B][C]
Generative reframing	Detecting that the criteria or boundary are wrong	No interface exists for purchasing or routing this capability. [A][B][E]

The third mode is absent not because individuals lack it, but because the system has no structural path for using it.

SECTION 9 — COMPENSATION AND HIDDEN BURDEN

The platform appears to function because a set of informal human carriers performs the continuity and interpretation functions the architecture does not formally support [B][C][E].

Queue Managers

Queue Managers (Attempter → Reviewer → Oracle → Queue Manager) are contributors, not employees [B]. Their documented functions — webinars, war rooms, Slack moderation — mask

their actual structural role: absorbing ambiguity. When guidelines change without notice, QMs translate the change. When scoring is opaque, QMs provide informal interpretation. When contributors cannot understand why their queue is empty, QMs offer the closest available explanation [C].

This is load-bearing coherence work. QMs cannot correct the frame — they can answer "how do I comply?" but not "is this guideline correctly framed?" The structural problem generating the ambiguity they absorb remains invisible to the only human layer with system-adjacent access [C][E].

Contributors themselves
Contributors carry significant unpaid coherence burden: repeated project-specific onboarding, queue-state inference assembled from community forums and third-party tools rather than platform information, and guideline interpretation effort not compensated as task time [C][D].

The existence of a community-built browser extension created specifically to surface hidden queue-state logic is evidence that contributors are compensating for missing platform legibility at the level of third-party tooling [C].

Structural implication: When informal carriers withdraw — when the contributor cannot reconstruct queue logic, when the QM burns out — the structural failure becomes immediately more visible. The system appears functional because these carriers absorb the ambiguity it produces. Their labour does not appear in any platform budget, productivity calculation, or correction proposal.

SECTION 10 — GOVERNING DIAGNOSIS

Outlier's recurring friction is not best explained as a contributor capability problem, a guideline clarity problem, or a routing inefficiency. The structural reading — derived from convergent public evidence across the recruitment surface, operating mechanics, symptom pattern, and consequence layer [A][B][C][D][E] — is this: the platform recruits through the legitimacy grammar of expertise while operating through the control grammar of a reliability pipeline. It signals that it is purchasing specialist judgment while its visible execution layer is designed to qualify, constrain, route, score, and repeat bounded task performance under variable demand. Because the intake architecture and the operating architecture describe different inputs, the platform predictably generates expectation mismatch at the qualification gate, routing opacity at the queue, compensatory support burden at the QM layer, quality variance at the scoring layer, and legal and reputational pressure at the consequence surface. Interface-level corrections cannot resolve this because the failure originates upstream of every interface — at the primitive. The architecture derives legitimacy benefits from expert-coded recruitment while operationally consuming compliance-governed output. These have never been the same thing, and no

amount of interface improvement will make them converge until the primitive is named and corrected.

The primitive mismatch — stated precisely:

	Stated	Actual
Primitive	Expert human judgment improves advanced AI systems	Variance-controlled instruction adherence at scale
What is purchased	Open-ended specialist reasoning	Repeatable rubric-conformant output
What is measured	Depth of domain judgment	Formatting, time adherence, rubric alignment
Who succeeds	Domain experts	Rubric-compliant performers
Who churns	High-signal specialists	—

Why the mismatch is structural, not incidental:
Reliability requires variance suppression. Genuine expertise requires that correct variance be valued. The platform cannot satisfy both requirements inside the same system with the same gate and the same scoring layer [A][B][E]. This is the structural contradiction.

Validation test:
If Outlier were to implement the Vanish List in Section 12 and internally declare its production primitive as "reliability pipeline with an optional expert judgment stream," would the remaining friction pattern be explained by execution quality rather than structural mismatch? If yes, the diagnosis holds. If specialist contributor churn persists even after the primitive is correctly named and the gate is rebuilt, revisit whether the fixed constraints in Section 2 prevent full primitive correction.

SECTION 11 — COST OF CONTINUATION

The following is a directional assessment of costs generated by the unresolved primitive mismatch. These are structural inferences from public evidence, not internal financial data.

Class 1 — Directly evidenced institutional cost

Legal and regulatory exposure [D]
A pattern of litigation and regulatory action consistent with the structural gap between the

platform's expert-facing promise and its reliability-pipeline operating reality. Three class action and PAGA lawsuits, a federal DOL investigation, and escalating investigative press coverage represent tail-risk exposure whose timing of crystallisation is uncertain.

Failed corrective infrastructure [A][C]
Additional compliance layers, support infrastructure, and interface redesigns have been deployed without primitive correction. Each produces incremental cost without structural resolution. The cost of the next corrective cycle is already embedded in the current failure pattern.

Class 2 — Structurally associated downstream burden

Recruitment waste [A][B][C]
The platform recruits, screens, verifies, and onboards contributors whose core capability the production layer is not primarily designed to consume. Every high-credential contributor who churns within one project cycle represents full intake cost with zero retained production value.

QA inflation [E]
Each compliance layer added in response to quality variance increases QA cost without reducing the variance, because the variance originates at the primitive, not at the execution layer.

Routing and support drag [B][C]
Queue opacity and project volatility generate permanent interpretive burden and QM support dependence. This cost cannot be reduced without correcting the underlying ambiguity that creates demand for it.

Class 3 — Invisible labour and burden transfer

Contributor-carried coherence labour [C][D]
The platform externalises coordination and interpretation cost onto contributors through unpaid onboarding burden, withheld queue logic, and community-dependent legibility. This labour does not appear in any platform budget or productivity calculation. It is the invisible input that keeps the system appearing to function.

The compounding dynamic:
Recruitment waste increases cost pressure → more algorithmic quality management → more scoring opacity → more contributor frustration → more public complaints → more compliance infrastructure → more signal-switch intensity → more churn among high-signal contributors → more recruitment waste. This loop has no internal exit without primitive correction [E].

Structural upside if corrected:
A correctly scoped reliability pipeline — with honest recruitment, a gate that selects for actual production need, and transparent queue logic — would materially reduce recruitment waste, QM overhead, legal exposure, and contributor churn simultaneously. Primitive correction should be understood as a recovery mechanism, not a discretionary improvement initiative.

SECTION 12 — CORRECTION DEPENDENCY ORDER

This is a dependency order, not a redesign specification. Right fixes in the wrong sequence still fail.

Vanish List — what stops before correction can hold:

Stop	Why
Using credential density as a proxy for production quality	Credentials predict expertise. The system is not primarily designed to consume expertise. The gate selects the wrong population. [A][B][E]
Treating longer guideline documents as a response to quality variance	More pages intensify the signal-switch. They convert expert labour to compliance labour earlier and more completely. [C][E]
Treating the Empty Queue as a communication problem	Internal EQ reason codes exist. The information exists. Withholding it is a design choice, not a limitation. [B][C]
Continued volume growth before primitive correction	Every additional credentialed contributor added before the primitive is corrected deepens the mismatch.

Move 1 — prerequisite for everything else:
Name the primitive correctly. The system must decide what it is actually purchasing: reliable repeatable instruction adherence at scale, genuine open-ended specialist judgment, or both — separated into structurally distinct streams with different gates, success criteria, and consequence paths. Without this declaration, every downstream fix will be calibrated to the wrong object [E].

Move 2 — follows from Move 1:
Redesign the qualification gate to select for the actual production need. The current gate selects for expertise signals. The platform is not primarily structured to consume expertise. The gate must be rebuilt to measure what the system actually requires. The redesign mechanism is in the full Redesign Artifact.

Move 3 — follows from Move 2:
Separate the two legitimate work classes currently collapsed into one broken frame. Collapsed

work classes must be structurally separated before correction can hold. The separation architecture is in the full Redesign Artifact.

Move 4 — follows from Move 3:
Install an upstream signal channel. No mechanism currently exists for a contributor who identifies a structural problem to surface that signal to a layer where it can be acted on. The channel design is in the full Redesign Artifact.

Move 5 — follows from Move 4:
Rebuild cross-project memory. The system currently starts every contributor's project assessment from zero. Demonstrated performance does not carry across project cycles. The memory architecture is in the full Redesign Artifact.

Dependency summary: Move 5 cannot precede Move 4. Move 4 cannot precede Move 3. Move 3 cannot precede Move 2. Move 2 cannot precede Move 1. The current pattern of interface corrections fails because it attempts Move 3 and Move 4 interventions without Move 1 in place.

SECTION 13 — PUBLIC EXEMPLAR NOTE

This document is a public exemplar derived from a full internal diagnostic artifact.

It is shared to demonstrate method, structural capability, and the distinction between diagnosis and redesign. The full internal artifact contains the deeper correction bridge, detailed transition logic, and builder handoff structure not disclosed here.

Full redesign specifications are not publicly available.

SECTION 14 — APPENDIX: EVIDENCE REFERENCES

Evidence tier key:
[A] First-party claims — Outlier/Scale AI official surfaces
[B] First-party operating signals — platform mechanics, governance documents
[C] Public symptom evidence — convergent contributor testimony and third-party tooling
[D] External consequence evidence — litigation, regulatory action, investigative press
[E] Structural inference sources — frameworks from which the diagnostic logic is derived

The following is a representative selection. Full source listing is in the master artifact.

[A][B] Official and operating signal sources (selected)

- Outlier.ai official materials: homepage, FAQ, specialist job pages, Marketplace announcements, Community Guidelines, Terms of Service
- Scale AI official materials: blog, public statements on contributor model
- Oxford Economics commissioned report (December 2025): 84% of contributors hold advanced degrees
- Contributor tier structure, pay rates, onboarding documentation

---

[C] Public symptom evidence (selected)

- Glassdoor, Indeed, Trustpilot: assessment opacity, guideline inconsistency, pay fluctuation, inability to appeal
- Reddit (r/outlier_ai, 3,600+ members): EQ patterns, onboarding loops, deactivation confusion
- BBB complaints (November 2024): unpaid onboarding hours, unanswered disputes
- GitHub: community-built browser extension surfacing internal EQ reason codes withheld from official interface

---

[D] External consequence evidence (selected)

- McKinney v. Scale AI (December 2024): class action, wage theft, misclassification, algorithmic pay reduction
- Amber Rogowicz PAGA suit (January 2025): effective pay below CA minimum wage
- Bloomberg Law (October 2024): WARN Act reporting
- TechCrunch (March 2025): U.S. DOL FLSA investigation
- The Guardian (April 2026): investigative reporting on task content and expert-facing promise divergence
- Oxford Internet Institute (2023): Remotasks (now Outlier) scored 1/10 on labour standards

---

[E] Structural inference sources

- *The Expertise Illusion in AI Task Marketplaces* (Forrester, February 2026): Transfer Test, interface-legitimacy mismatch, missing coordination layer
- *The Four-Function Law of Scalable Institutions* (Forrester, March 2026): function fusion, sensing-authority collapse, fragmented interpretation, absent cross-project memory

---

This public diagnostic exemplar was produced by Jamie Forrester, Independent Systems Architect, Edinburgh, UK. It is based entirely on publicly available information and was produced independently — it was not commissioned by Outlier or Scale AI. The diagnostic logic, structural interpretation, and governing diagnosis are the author's own.

This public exemplar is shared to demonstrate method and structural capability. Full internal artifacts and redesign specifications are not publicly available.

hello@jamieforrester.com | jamieforrester.com